



Efficient Execution Plans for Distributed Skyline Query Processing

Paper by

Joao B. Rocha Junior, Akrivi Vlachou

Christos Dolkeridis, Kjetil Norvag

Presentation by

Ritika (15111036)

Swapnil Mhamane(15111044)

Published in EDBT/ICDT '11

System Overview

- Each server stores autonomously a fraction of the data.
 - Horizontal partition of data.
- All servers need to process the skyline query.
- Each server S_i can directly connect to any other server S_j .

Naïve Approach

- A skyline query can be initiated by any server (S_{org}).
- Skyline query is processed by sending the query to all servers S_i .
- Each server S_i reports its local skyline set SKY_i to originator.
- S_{org} gathers local skyline set and computes global skyline set
- Use of transitivity property
- Single hop execution

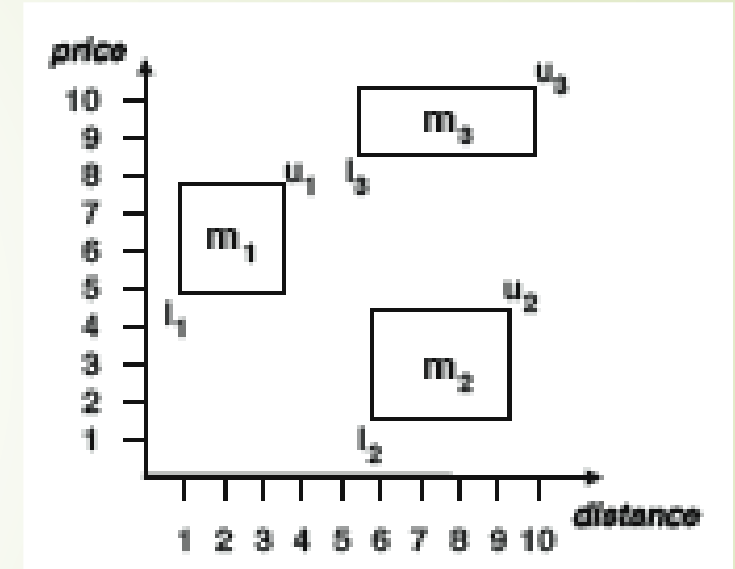
Motivation

- Execution plan defines the order in which the individual skyline queries are processed on different servers
- Servers need not be contacted can be pruned
- Discarding some data points locally based on skyline points from preceding server
- Negative point : Local skyline queries on consecutive server is a blocking operation
- Observation :
 - There exist **dependency** between two servers
 - Optimizations by exploiting these dependencies.

Relation among MBR

- A hyper-rectangle $m_i (l_i, u_i)$.
- Given two hyper-rectangles m_i and m_j
 - (1) m_i dominates m_j , if $u_i < l_j$;
 - (2) m_i partially dominates m_j , if $l_i < u_j$, but $u_i \nless l_j$;
 - (3) m_i and m_j are incomparable, if $l_i \nless u_j$ and $l_j \nless u_i$.
- Enclosed dominance area V_{ij} of m_i on m_j : the volume of m_j that is dominated by the lower left corner l_i of m_i
- **Pruning power** (PP_{ij}) of m_i on m_j :

$$PP_{ij} = V_{ij}/V_j$$



Main phases of SkyPlan


- Collect the MBRs of all servers.
- Builds a weighted directed graph
- The graph is transformed into an execution plan
- Execute plan recursively

Skyline dependency graph


- ▶ SD-graph : the weighted directed graph $G(N, E, w)$
- ▶ Each node $n_i \in N$ corresponds to a non-dominated MBR m_i .
- ▶ E is a set of ordered pairs $e_{ij} = (i, j)$, where m_i, m_j are MBRs, and m_i partially dominates m_j
- ▶ w is a weight function :
 - ▶ Normalized pruning power

$$w_{ij} = |m_j| / |D| * PP_{ij}$$

SD-graph example




(a) Set of MBRs.




(b) SD-graph.


Finding Maximum Spanning Tree




(a) MBRs.




(b) SD-graph.



(c) Dummy node.



(d) Spanning tree.



(e) Plan.

Execution Plan

- The quality $Q(P)$ of an execution plan $P(N,E)$

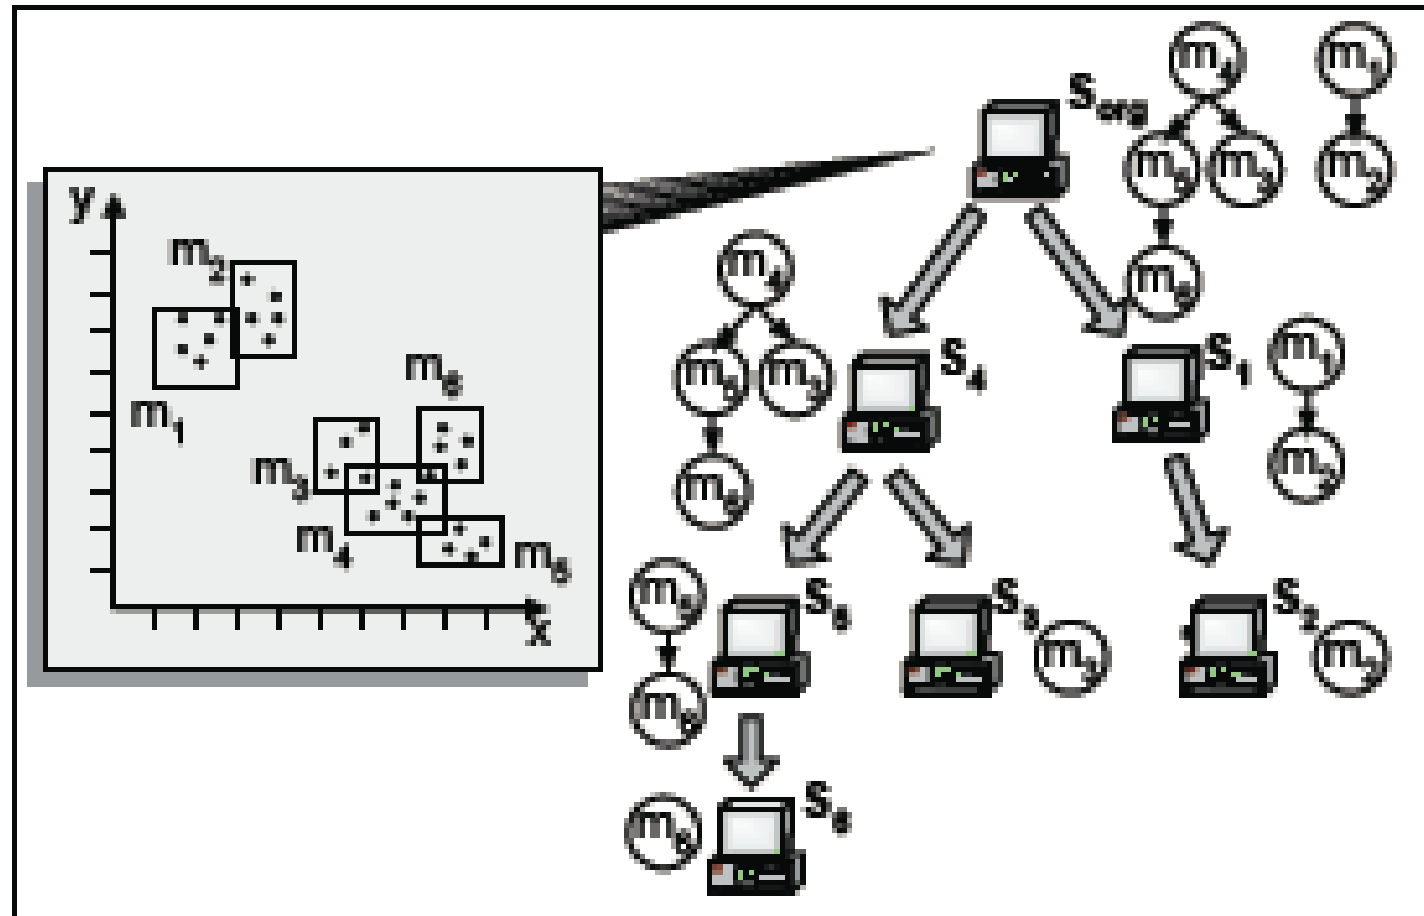
$$Q(P) = \sum_{e_{ij} \in E} w_{ij}$$

- Thus, an execution plan P_i is better than another plan P_j in terms of quality if: $Q(P_i) > Q(P_j)$.
- Maximum pruning execution plan
- Execution plan : Set of directed weighted trees (Forest)
- Each tree is processed in parallel

Algorithm: QueryProcessing(S_i, F, P)

- **INPUT:** Filter points $F = \{f_1, \dots, f_k\}$, Execution plan P .
- **OUTPUT:** Local skyline
- $m \leftarrow P.\text{getRootMBR}()$
- $\text{sky} \leftarrow \text{computeSkyline}(m)$
- $P' \leftarrow \text{refinePlan}(P, \text{sky})$
- $F' \leftarrow \text{refineFilters}(F, \text{sky})$
- $S' \leftarrow P'.\text{getNextServers}()$
- for ($\forall S_j \in S'$) do $\text{sky}_j \leftarrow \text{QueryProcessing}(S_j, F', P'_j)$
- $\text{sky} \leftarrow \text{mergeSkyline}(\text{sky}, \text{sky}_j)$
- return sky

Example



K-hop execution plan

- Another objective is to restrict the number of hops in query processing
 - Bounding latency in consecutive execution
- K-hop execution plan : Execution plan with height at most k
- Hop constrained maximum spanning tree problem
- **Algorithm :**
 - Find longest path in generated execution plan
 - Replace pruning power edge with other edge reducing length by at least 1
 - If there exist more than one such edges, choose edge with maximum pruning power
 - Repeat until length is reduced to k

Summary

- Paper targets the problem of deriving efficient execution plan for distributed skyline computation
- It proposes the novel framework called SkyPlan, that maps the dependencies between the queries into graph and generates cost aware execution plan.
- Aim was to maximize pruning power in consecutive queries while keep increment in parallelism
- It proposes distributed query execution mechanism that allows continuous refinement of plan during in-network query processing.

References

- João B. Rocha-Junior* , Akrivi Vlachou, Christos Doulkeridis, and Kjetil Nørvåg “Efficient execution plans for distributed skyline query processing”

Thank you ...!

